

Visuomotor Cooperative Manipulation with Teams of Quadrotors

Cooperative quadrotor teaming has huge potential in repetitive tasks in construction, agriculture, or people and goods transport where lifting and transporting objects is necessary. Recent methods have shown that it is possible to cooperatively lift and transport payloads using teams of drones, but these methods rely on motion capture systems, explicit knowledge of the payload/object and fixed attachment points between the payload and the drones. In this project we will be enabling teams of two or three quadrotors to cooperatively lift and manipulate the pose of arbitrary objects using only onboard sensing and perception using recent advancements in robot learning such as large-scale pretrained diffusion policies and continual learning methods using reinforcement learning and techniques from control theory. As a first step the student who will be selected will perfect manipulation of a single object using a quadrotor with a single gripper attached using only onboard sensing. The research will be conducted at the University of Cambridge in the Division of Information Engineering (Machine Intelligence Group, Mandralis Lab). Visiting students from ETHZ or EPFL will benefit from a highly stimulating intellectual environment and tradition at the University of Cambridge.

PROJECT	Master's thesis	STUDENT	TBD
SEMESTER	Jan–July 2027	HOST	University of Cambridge
START	January 2027	SUPERVISOR	Prof. Ioannis Mandralis
DEADLINE	July 2027	ADVISORS	—
CONTACT	imm38@cam.ac.uk		

Keywords aerial manipulation; cooperative control; visuomotor policies; diffusion policies; reinforcement learning; onboard perception

1 Introduction

Aerial robots that can grasp, lift and reposition objects would extend automation to tasks in construction, agriculture and logistics that ground robots cannot reach, and where the payload often exceeds what a single vehicle can carry [1, 2]. Teams of quadrotors can share such loads, and cooperative transport has been demonstrated in several forms, from multiple vehicles rigidly grasping a common object [3] to teams carrying cable-suspended payloads, including in agile flight [4, 5].

These demonstrations rely on assumptions that do not hold outside the laboratory. The states of the vehicles and of the payload are usually provided by a motion capture system; the mass, inertia and

geometry of the payload are known in advance; and the payload is attached to the vehicles at fixed, pre-defined points. Work towards onboard sensing exists [6, 7], but it addresses the transport of a known payload rather than the grasping and reorientation of an object whose properties are not known beforehand.

In parallel, visuomotor policies trained on large demonstration datasets, such as diffusion policies and vision-language-action models, allow fixed-base manipulators to grasp and manipulate a wide range of objects from camera images alone [8, 9, 10]. Recent work has begun to transfer such policies to aerial manipulators by accounting for the dynamics of the flying platform when actions are generated [11]. Reinforcement learning in simulation, followed by transfer to hardware, has produced quadrotor controllers that operate close to the physical limits of the vehicle [12, 13].

This thesis combines these directions. The student will first develop a single quadrotor with a gripper that grasps, lifts and reorients objects using only onboard sensing: a pretrained visuomotor policy provides end-effector targets from camera images, and a low-level controller, trained with reinforcement learning or derived from control theory, tracks them while compensating for the unknown payload. The approach will then be extended to teams of two or three quadrotors that jointly grasp an object and control its pose, selecting their own grasp points and coordinating through onboard perception rather than external tracking.

2 Objectives

This project aims to:

1. **Onboard platform.** Build a quadrotor with a gripper and an onboard sensing and compute stack (visual-inertial odometry, an RGB-D camera and an embedded GPU) that operates without motion capture. The lab's existing quadrotors will serve as the starting point.
2. **Simulation and data.** Develop a simulation environment for one or more quadrotors with grippers interacting with objects of varying shape, mass and inertia, suitable for massively parallel reinforcement learning and for generating demonstrations with a privileged expert controller.
3. **Single-quadrotor visuomotor manipulation.** Fine-tune a large-scale pretrained visuomotor policy (e.g. a diffusion policy or a vision-language-action model) to grasp, lift and reorient previously unseen objects with a single quadrotor, and refine the closed-loop behaviour through reinforcement learning in simulation and continual learning on hardware.
4. **Cooperative manipulation.** Extend the approach to two or three quadrotors that grasp a common object at self-selected points and control its position and orientation, using tools from control theory to regulate the internal forces between vehicles and to guarantee stable behaviour of the coupled system.
5. **Experimental validation.** Quantify the performance of the onboard system against a baseline that uses motion capture and known payload parameters, in terms of pose-tracking error, success rate over a set of unseen objects, and robustness to disturbances.

3 Work breakdown

The project will be structured into the following work packages:

WP 1 Literature review and platform

Review work on cooperative aerial manipulation, visuomotor policy learning and learning-based quadrotor control. Assemble and characterise the quadrotor-gripper platform with its onboard sensing and compute, and validate state estimation against motion capture.

WP 2 Simulation and demonstration data

Implement the quadrotor–gripper–object system in a GPU-accelerated simulator with contact, and randomise object geometry, mass, inertia and friction. Generate demonstrations with a privileged expert (e.g. model predictive control with full state knowledge) and, where useful, by teleoperation on hardware.

WP 3 Single-quadrotor visuomotor grasping

Fine-tune a pretrained visuomotor policy on the demonstration data so that it outputs end-effector targets from onboard images. Interface it with a low-level controller that tracks these targets, and demonstrate grasping, lifting and reorientation of a single object on hardware.

WP 4 Policy refinement


Train the low-level controller, or a residual on top of the pretrained policy, with reinforcement learning in simulation to compensate for unknown payloads. Investigate continual learning on hardware so that the policy improves from the robot's own experience.

WP 5 Cooperative manipulation with two or three quadrotors

Extend the system to multiple vehicles grasping one object. Study decentralised policies that coordinate through onboard perception and limited communication, and use control-theoretic tools to regulate internal forces and ensure stability of the coupled system.

WP 6 Experimental evaluation

Run an experimental campaign on a set of objects not seen in training and compare against a motion-capture baseline with known payload parameters. Report pose-tracking error, success rate and robustness to disturbances.

3.1 Indicative timeline**4 Candidate profile**

We are looking for a student with:

- A background in robotics, control or machine learning, from mechanical, electrical or computer engineering, computer science or a related programme.
- Programming experience in Python and PyTorch; experience with C++ and ROS 2 is an advantage.
- Familiarity with reinforcement learning, imitation learning or model-based control.
- Interest in experimental work; hands-on experience with drones or embedded hardware is an

advantage.

To apply, send a CV, a transcript of records and a short statement of interest to the contact address above.

5 Deliverables and assessment

The student is expected to deliver:

1. A written report following the host university's guidelines, submitted by the deadline stated above.
2. A final presentation to the laboratory (20 min plus questions).
3. A documented code repository, comprising the simulation environment, the policy training and controller implementations, and the experimental data, structured so the work can be reproduced and carried on in follow-up projects.

Assessment follows the host university's regulations, with weight on the scientific quality of the work, the independence shown during the project, the quality of the documentation and the reproducibility of the results.

6 References

- [1] A. Ollero et al., "Past, Present, and Future of Aerial Robotic Manipulators," *IEEE Transactions on Robotics*, 2022.
- [2] K. Zhang et al., "Aerial additive manufacturing with multiple autonomous robots," *Nature*, 2022.
- [3] D. Mellinger et al., "Cooperative Grasping and Transport Using Multiple Quadrotors," *Distributed Autonomous Robotic Systems*, 2013.
- [4] M. Tognon et al., "Aerial Co-Manipulation With Cables: The Role of Internal Force for Equilibria, Stability, and Passivity," *IEEE Robotics and Automation Letters*, 2018.
- [5] S. Sun et al., "Agile and cooperative aerial manipulation of a cable-suspended load," *Science Robotics*, 2025.
- [6] G. Loianno and V. Kumar, "Cooperative Transportation Using Small Quadrotors Using Monocular Vision and Inertial Sensing," *IEEE Robotics and Automation Letters*, 2018.
- [7] G. Li, R. Ge, G. Loianno, "Cooperative Transport of a Suspended Payload via Two Aerial Robots With Inertial Sensing," *IEEE Access*, 2021.
- [8] C. Chi et al., "Diffusion Policy: Visuomotor Policy Learning via Action Diffusion," *Robotics: Science and Systems*, 2023.
- [9] C. Chi et al., "Universal Manipulation Interface: In-The-Wild Robot Teaching Without In-The-Wild Robots," *Robotics: Science and Systems*, 2024.
- [10] K. Black et al., " π_0 : A Vision-Language-Action Flow Model for General Robot Control," *arXiv preprint arXiv:2410.24164*, 2024.
- [11] H. Gupta et al., "UMI-on-Air: Embodiment-Aware Guidance for Embodiment-Agnostic Visuomotor Policies," *arXiv preprint arXiv:2510.02614*, 2025.
- [12] J. Hwangbo et al., "Control of a Quadrotor With Reinforcement Learning," *IEEE Robotics and Automation Letters*, 2017.
- [13] E. Kaufmann et al., "Champion-level drone racing using deep reinforcement learning," *Nature*, 2023.

Cambridge, TBD

SUPERVISOR

STUDENT

Prof. Ioannis Mandralis

University of Cambridge